

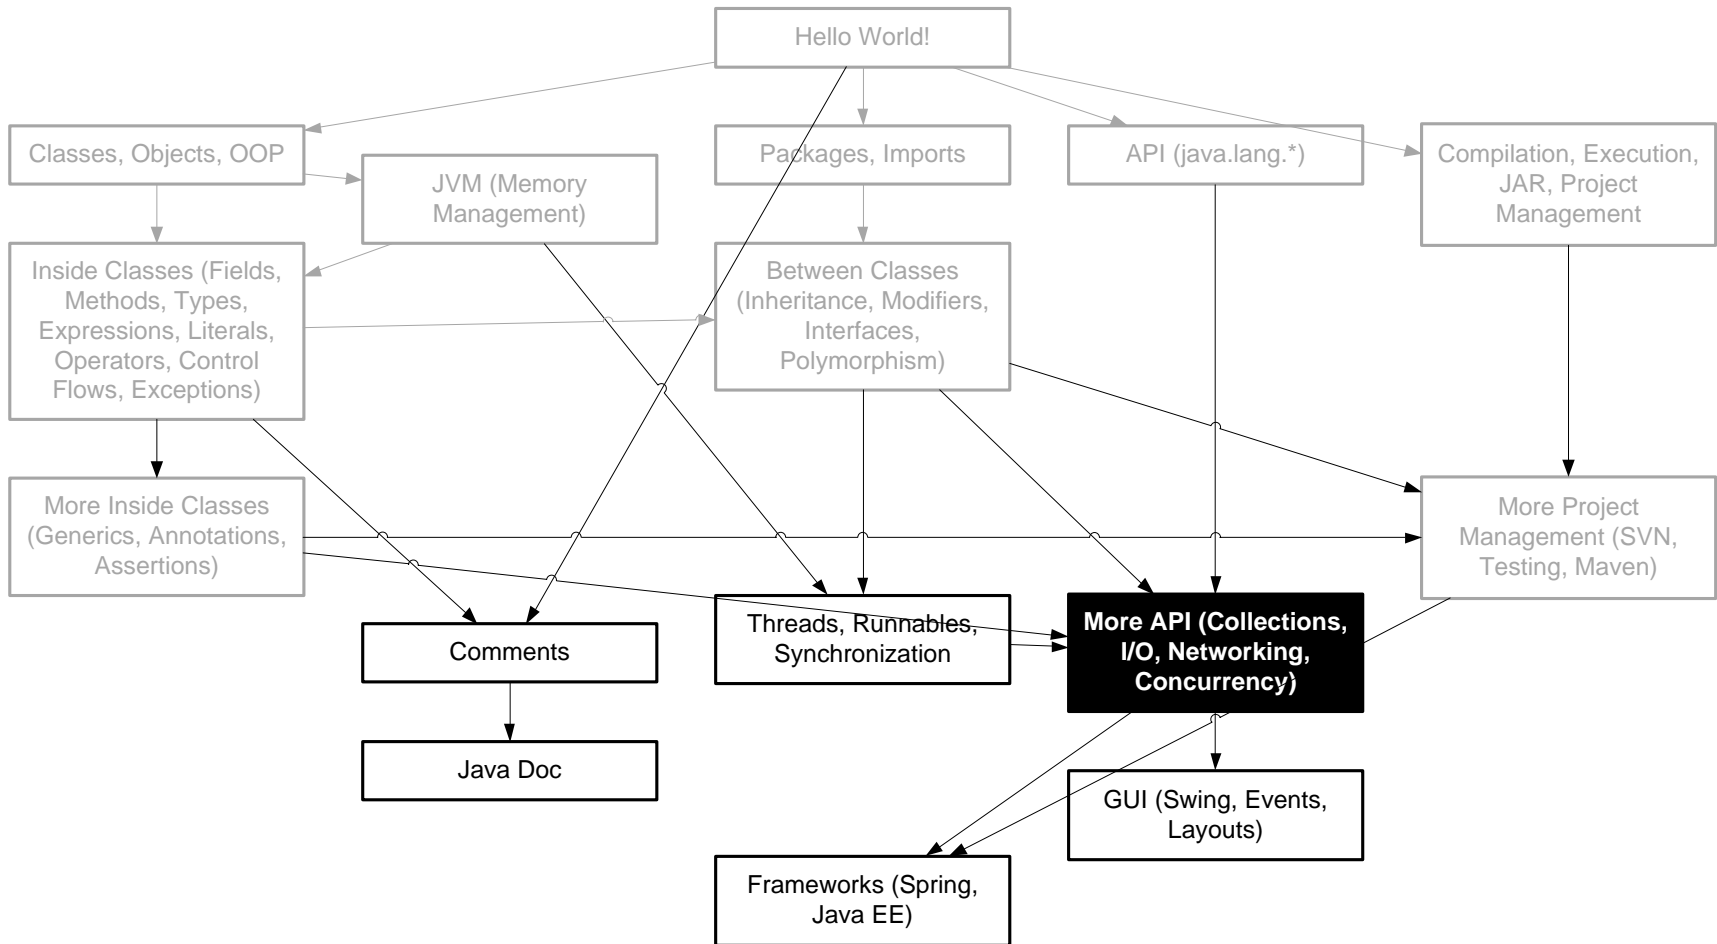
Java I/O

Shan-Hung Wu

CS, NTHU,

Fall, 2013

Where Are We?



Why I/O?

- I/O is one of the most common tasks for every programmer
- Main packages:
 - `java.io`: traditional API that provides high-level abstractions for paths, random access files, and streams
 - `java.nio`: complement `java.io` with low-level abstractions and asynchronous access

Assigned Readings

- For each class/method mentioned in this lecture, read the corresponding document in Java API spec

Agenda

- Paths, Files, and Directories
 - java.io.File
 - java.io.RandomAccessFile
- Streams
 - Byte Streams
 - Buffered Streams
 - Character Streams and Formatting
 - Data Streams
 - Object Streams
- New I/O
 - Buffers and Channels
 - More File I/O

Paths

- `java.io.File` provides a system-independent view of hierarchical **paths** (which may not exist)
- Represents either files or directories:
 1. `File parent = new File("/bin");` // a directory
 2. `File child = new File(parent, "java.exe");` // a file
 - '/' works on all platforms, including Windows
 - Paths are relative to the working directory if not start with '/'
- Usage:
 - Check for existence and permissions: `exist()`, `canWrite()` etc.
 - Query various info: `isFile()`, `length()`, `lastModified()`, etc.
 - Create, rename or delete files and directories: `createNewFile()`, `createTempFile()`, `mkdir()`, `renameTo()`, `delete()`, `deleteOnExit()`, etc.
- Static fields `File.separator` and `File.pathSeparator` provide quick access to system-dependent separators

Tips for Path Manipulation

- To obtain a temporary file:

```
1. File temp = File.createTempFile("temp-", ".jpg", "/dir");  
2. temp.deleteOnExit(); // optional
```

- Each temp file will be guaranteed to have a unique name, e.g., temp-256408564856505064.jpg

- Enumerate Windows/Unix roots:

```
1. File[] winDrives = File.listRoots(); // for win  
2. File unixRoot = File.listRoots()[0]; // for unix
```

- Enumerate files in a directory:

```
1. File[] files = new File("some dir").listFiles();
```

Random Access Files

- `java.io.RandomAccessFile` provides an abstraction to random accessible files (usually on local disks)
 1. `File path = ...`
 2. `RandomAccessFile rf = new RandomAccessFile(path, "rws");`
 3. `rf.seek(10);` // set pointer at which next read/write occurs
 4. `int b = rf.read();` // pointer is advanced automatically
 5. `rf.seek(10);`
 6. `rf.write(++b);`
- Mode: "r", "rw", "rws"
 - 's' ensures that every write is synchronized to disk
- Acts like a large array of bytes
 - Has a **file pointer** that can be read and moved:
`getFilePointer()`, `seek()`

RandomAccessCopier

```
1. public class RandomAccessCopier implements FileCopier {
2.     @Override
3.     public void copy(File src, File dest) throws IOException {
4.         RandomAccessFile srcRa = null, destRa = null;
5.         try {
6.             srcRa = new RandomAccessFile(src, "r");
7.             destRa = new RandomAccessFile(dest, "rws");
8.             int b = -1;
9.             while ((b = srcRa.read()) != -1) {
10.                 destRa.write(b);
11.             }
12.         } finally {
13.             if (srcRa != null) srcRa.close();
14.             if (destRa != null) destRa.close();
15.         }
16.     }
17. }
```

Performance Issues

- `destRa` (with mode "rws") ensures that each call to `write()` is reflected to disk
- However, disks are very slow
 - Typically, the access time of RAM is about 60 ns
 - 1,000 times faster than flash drives (~60 us)
 - 100,000 times faster than magnetic drives (~6 ms)
- The performance of `RandomAccessCopier` is poor
- Can you make it faster?

BufferedRandomAccessCopier

```
1. public class BufferedRandomAccessCopier implements FileCopier {
2.     private static final int BUFFER_SIZE = 8192; // in bytes
3.
4.     @Override
5.     public void copy(File src, File dest) throws IOException {
6.         RandomAccessFile srcRa = null, destRa = null;
7.         byte[] buffer = new byte[BUFFER_SIZE];
8.         try {
9.             srcRa = new RandomAccessFile(src, "r");
10.            destRa = new RandomAccessFile(dest, "rws");
11.            int num = -1;
12.            while ((num = srcRa.read(buffer)) != -1) {
13.                destRa.write(buffer, 0, num);
14.            }
15.        } finally {
16.            if (srcRa != null) srcRa.close();
17.            if (destRa != null) destRa.close();
18.        }
19.    }
20.}
```

Agenda


- Paths, Files, and Directories
 - `java.io.File`
 - `java.io.RandomAccessFile`
- Streams
 - Byte Streams
 - Buffered Streams
 - Character Streams and Formatting
 - Data Streams
 - Object Streams
- New I/O
 - Buffers and Channels
 - More File I/O

Streams (1/2)

- `File` and `RandomAccessFile` corresponding to paths and files in your file system
 - They are API specific to **file I/O**
- **General I/O** can read/write bytes from/to any device other than disks
 - E.g., network interface card, console, etc.
- Random access may not be possible!

Streams (2/2)

- `java.io.InputStream/OutputStream` provides an abstraction of **streams**



- Streams are independent of the nature of data sources
 - Bytes can be read/written sequentially by calling `read()` and `write()` repeatedly

ByteStreamCopier

```
1. public class ByteStreamCopier implements FileCopier {
2.     @Override
3.     public void copy(File src, File dest) throws IOException {
4.         InputStream in = null;
5.         OutputStream out = null;
6.         try {
7.             in = new FileInputStream(src);
8.             out = new FileOutputStream(dest);
9.             int b;
10.            while ((b = in.read()) != -1) {
11.                out.write(b);
12.            }
13.        } finally {
14.            if (in != null) in.close();
15.            if (out != null) out.close();
16.        }
17.    }
18.}
```

Types of Streams

- Two main categories in Java:
 - Byte streams: `InputStream/OutputStream`
 - Character streams: `Reader/Writer`
- Endpoint streams that read/write directly from/to devices:
 - E.g., `FileInputStream, FileReader, ByteArrayInputStream, StringReader, etc.`
- Wrapper streams that read/write from/to endpoint streams:
 - E.g., `BufferedOutputStream, BufferedWriter, DataOutputStream, ObjectOutputStream, PrintWriter etc.`

BufferedByteStreamCopier

```
1. public class BufferedByteStreamCopier implements FileCopier {
2.     @Override
3.     public void copy(File src, File dest) throws IOException {
4.         InputStream in = null;
5.         OutputStream out = null;
6.         try {
7.             in = new BufferedInputStream(new FileInputStream(src));
8.             out = new BufferedOutputStream(new FileOutputStream(dest));
9.             int b;
10.            while ((b = in.read()) != -1) {
11.                out.write(b);
12.            }
13.            // (optional) make sure data are reflected to disk
14.            out.flush();
15.            ...
16.        } finally {
17.            if (in != null) in.close();
18.            if (out != null) out.close(); // flush first
19.        }
20.    }
21.}
```

CharStreamCopier

```
1. public class CharStreamCopier implements FileCopier {
2.     private static final String CHAR_SET = "UTF-8";
3.     @Override
4.     public void copy(File src, File dest) throws IOException {
5.         Reader r = null;
6.         Writer w = null;
7.         try {
8.             r = new InputStreamReader(
9.                 new FileInputStream(src), CHAR_SET);
10.            w = new OutputStreamWriter(
11.                new FileOutputStream(dest), CHAR_SET);
12.            int c; // char as int
13.            while ((c = r.read()) != -1) {
14.                w.write(c);
15.            }
16.        } finally {
17.            if (r != null) r.close();
18.            if (w != null) w.close();
19.        }
```

- Java also provides FileReader/FileWriter, but they support default encoding only

LinedCharStreamCopier

```
1. public class LinedCharStreamCopier implements FileCopier {
2.     private static final String CHAR_SET = "UTF-8";
3.     @Override
4.     public void copy(File src, File dest) throws IOException {
5.         BufferedReader br = null;
6.         BufferedWriter bw = null;
7.         try {
8.             br = new BufferedReader(
9.                 new InputStreamReader(
10.                    new FileInputStream(src), CHAR_SET));
11.             bw = new BufferedWriter(
12.                 new OutputStreamWriter(
13.                    new FileOutputStream(dest), CHAR_SET));
14.             String s;
15.             while ((s = br.readLine()) != null) { // system independent
16.                 bw.write(s);
17.             }
18.         } finally {
19.             if (br != null) br.close();
20.             if (bw != null) bw.close();
21.         }
22.     }
23. }
```

Formatting


- Char streams simplify reading/writing lines
 - System independent; works either for `"\r"`, `"\r\n"`, or `"\n"`
- Another advantage: formatting
 1. `PrintWriter pw = new PrintWriter(/* some Writer */);`
 2. `// or`
 3. `PrintStream pw = new PrintStream(/* some OutputStream */,
..., /* charset */)'`
 4. `pw.format("The square root of %d is %f. Correct?%n",
9, Math.sqrt(9));`
 5. `...`
 6. `System.out.format("Current time at %s is %tc.%n",
"Taipei", System.currentTimeMillis());`
- Format specifier: `%`
- Conversion: `n` (line terminator), `d` (decimal), `f` (floating point), `s` (string), `tH` (hour), `tM` (minute), `tC` (date/time), etc.

Formatting (2/2)

```
1. System.out.format("<%f, %1$+020.10f>%n", Math.PI);  
   // output: <3.141593, +000000003.1415926536>
```

- Flags:

- + (signed)
- 0/ ' ' (0/space-padded)
- - (right padding)
- , (thousand separated)



- Width: minimum width; padded if necessary
- Precision: for floating points; truncated at right
- Read this [format syntax API spec](#) for more details

DataStream

- `java.io.DataInputStream/DataOutputStream` allows reading/writing primitive and String values from/to streams in a portable way
 - `writeBoolean()`, `writeInt()`,
`writeDouble()`, `writeChar()`,
`writeUTF()`, etc.

ObjectStream

- `java.io.DataInputStream/DataOutputStream` allows reading/writing primitives **and objects** from/to streams

```
1. ObjectOutputStream oos = new ObjectOutputStream(...);
2. oos.writeObject("Today");
3. oos.writeObject(new Date());
4. ClassA a1 = new ClassA(...);
5. oos.writeObject(a1);
6. oos.close();
7. ...
8. ObjectInputStream ois = new ObjectInputStream(...);
9. String s = (String) ois.readObject();
10. Date date = (Date) ois.readObject();
11. ClassA a2 = (ClassA) ois.readObject();
12. ois.close();
13....
14. System.out.println(a1 == a2); // true or false?
```

Object Serialization (1/2)

- We say objects are **deserialized/serialized** upon reading/writing from streams
 - Unlike cloning, serialization is deep; all fields will be serialized recursively
 - The only exceptions are those fields declared with the `transient` modifier
- `ObjectInputStream/OutputStream` provides a convenient way to perform deep cloning:

```
1. ObjectOutputStream oos = new ObjectOutputStream(  
    new ByteArrayOutputStream()); // endpoint is memory  
2. ClassA obj = new ClassA(...);  
3. oos.writeObject(obj);  
4. ObjectInputStream ois = new ObjectInputStream(  
    new ByteArrayInputStream(oos.toByteArray()));  
5. ClassA deepCloneObj = (ClassA) ois.readObject();  
6. ...  
7. oos.close();  
8. ois.close();
```


Object Serialization (2/2)

- An object is **not** serializable unless its class implements `java.io.Serializable`

```
1. public ClassA implements Serializable {  
2.     private static final long serialVersionUID = 1L;  
3.     ...  
4. }
```

- What happens if you serialize an object of ClassA, modify ClassA, and then deserialize the object?
 - A serializable class should define a field named `serialVersionUID`
 - Increment `serialVersionUID` if you think the modification of ClassA makes the previously serialized objects incompatible with the new ones

Agenda

- Paths, Files, and Directories
 - `java.io.File`
 - `java.io.RandomAccessFile`
- Streams
 - Byte Streams
 - Buffered Streams
 - Character Streams and Formatting
 - Data Streams
 - Object Streams
- New I/O
 - Buffers and Channels
 - More File I/O

Channels and Buffers

- `java.nio` offers a new way for I/O: Channels and Buffers
- All data operations are performed on Buffers (data blocks in memory)
 - There are buffers for all primitive types, like `ByteBuffer`, `LongBuffer`, `FloatBuffer`, etc.
 - Buffers can be allocated directly by OS, saving memory copying from user space to kernel space before each I/O operation
 - Buffers maintain their capacity, limit, and position, thus simplifying your code
- Channels are like streams, but they
 - Are bi-directional
 - Read/write data from/into Buffers only

Buffer Positions

- Buffers maintain their own positions
 - $0 \leq \text{mark} \leq \text{position} \leq \text{limit} \leq \text{capacity}$
- `clear()`: sets the limit to the capacity and the position to zero
 - Makes the buffer ready for channel-read
- `flip()`: sets the limit to the current position and then sets the position to zero
 - Makes the buffer ready for channel-write
- `rewind()`: leaves the limit unchanged and sets the position to zero
 - Makes the buffer ready for re-reading the data that it already contains

ChannelCopier

```
1. public class ChannelCopier implements FileCopier {
2.     private static final int BUFFER_SIZE = 8192; // in bytes
3.     private ByteBuffer buffer = ByteBuffer.allocateDirect(BUFFER_SIZE);
4.     @Override
5.     public void copy(File src, File dest) throws IOException {
6.         FileChannel ic = null, oc = null;
7.         try {
8.             ic = new FileInputStream(src).getChannel();
9.             oc = new FileOutputStream(dest).getChannel();
10.            while (ic.read(buffer) != -1) {
11.                buffer.flip();
12.                oc.write(buffer);
13.                buffer.clear();
14.            }
15.        } finally {
16.            if (ic != null) ic.close();
17.            if (oc != null) oc.close();
18.        }
19.    }
20.}
```

More File I/O

- In Java 7, `java.nio` is further extended to support advanced file I/O operations
 - Recursive file traversal
 - Manipulating symbolic/hard links
 - Searching files with wildcards, etc.
- Here is a [mapping](#) between the new and old file I/O APIs
- Optional reading: [File I/O in Java 7](#)