

Implementing Web Intelligence: An Analyzer Framework

NetDB

CS, NTHU,

Fall, 2013

Now you know how a recommendation engine work

- Strategy
 - User-based
 - Content-based
- How about building up a content-based recommendation engine by yourself!

Our Tiny Recommendation Engine

- To do recommendation, we need to
 1. Analyze(Parse) a set of documents based on "terms"
 2. Store the analyzed statistics about the "terms"
 3. Given a document and an integer k , searches and returns the k most similar documents in the store
- You can reuse this engine in
 - Your final project
 - Or even future projects

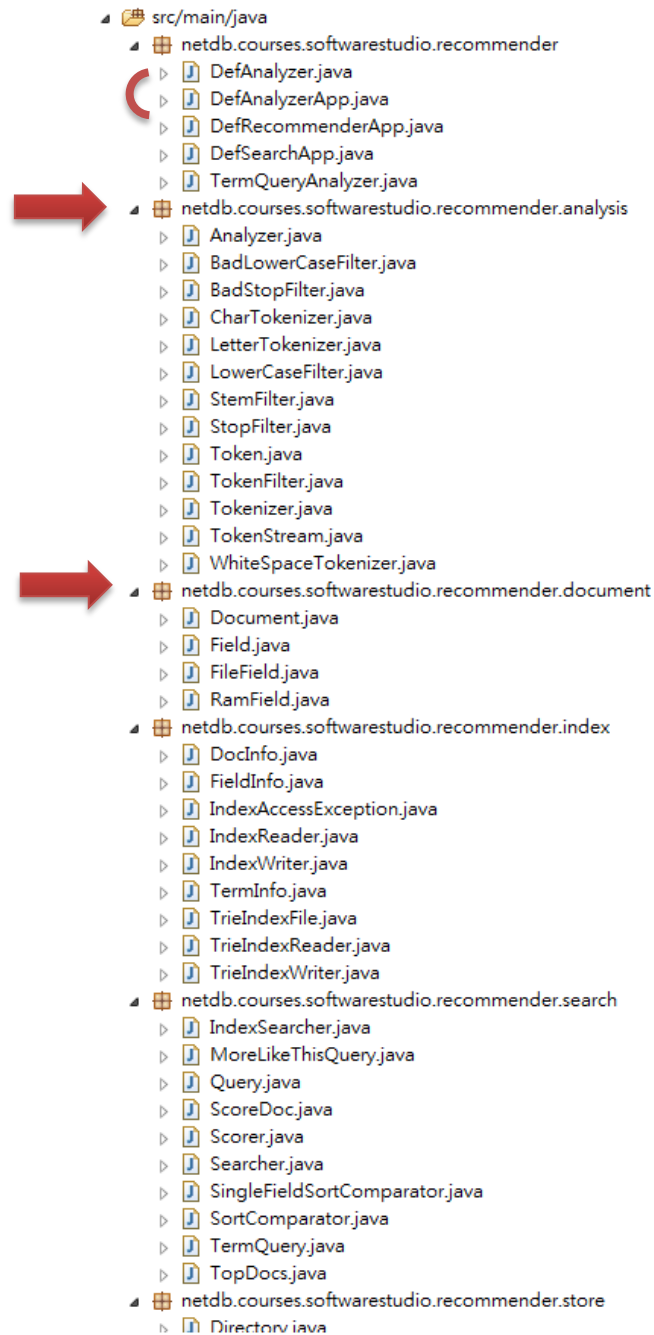
More Detailed Steps

1. Create a **Document** class that is general enough to contain the searchable content
 - E.g., the content of a Definition object, a WiKi page, etc.
2. Implement an **Analyzer** that extracts the “terms” from a Document object (doc for short)
3. An **IndexWriter** that calculates the statistics about the terms/docs (e.g., term frequencies normalized by doc norms) and stores them into an **Index**
 - Basically, an index is a map from the terms to document statistics and IDs
4. An **IndexSearcher** that to find the IDs of the most similar docs to a given doc from index
 - Calculate cosine similarity by looking up the index term-by-term

- src/main/java
 - netdb.courses.softwarestudio.recommender
 - DefAnalyzer.java
 - DefAnalyzerApp.java
 - DefRecommenderApp.java
 - DefSearchApp.java
 - TermQueryAnalyzer.java
 - netdb.courses.softwarestudio.recommender.analysis
 - Analyzer.java
 - BadLowerCaseFilter.java
 - BadStopFilter.java
 - CharTokenizer.java
 - LetterTokenizer.java
 - LowerCaseFilter.java
 - StemFilter.java
 - StopFilter.java
 - Token.java
 - TokenFilter.java
 - Tokenizer.java
 - TokenStream.java
 - WhiteSpaceTokenizer.java
 - netdb.courses.softwarestudio.recommender.document
 - Document.java
 - Field.java
 - FileField.java
 - RamField.java
 - netdb.courses.softwarestudio.recommender.index
 - DocInfo.java
 - FieldInfo.java
 - IndexAccessException.java
 - IndexReader.java
 - IndexWriter.java
 - TermInfo.java
 - TrieIndexFile.java
 - TrieIndexReader.java
 - TrieIndexWriter.java
 - netdb.courses.softwarestudio.recommender.search
 - IndexSearcher.java
 - MoreLikeThisQuery.java
 - Query.java
 - ScoreDoc.java
 - Scorer.java
 - Searcher.java
 - SingleFieldSortComparator.java
 - SortComparator.java
 - TermQuery.java
 - TopDocs.java
 - netdb.courses.softwarestudio.recommender.store
 - Directory.java

Today's Topic

1. Create a **Document** class that is general enough to contain the searchable content
 - E.g., the content of a Definition object, a WiKi page, etc.
2. Implement an **Analyzer** that extracts the “terms” from a Document object (doc for short)
3. An IndexWriter that calculates the statistics about the terms/docs (e.g., term frequencies normalized by doc norms) and stores them into an Index
 - Basically, an index is a map from the terms to document statistics and IDs
4. An IndexSearcher that to find the IDs of the most similar docs to a given doc from index
 - Calculate cosine similarity by looking up the index term-by-term



Code Time

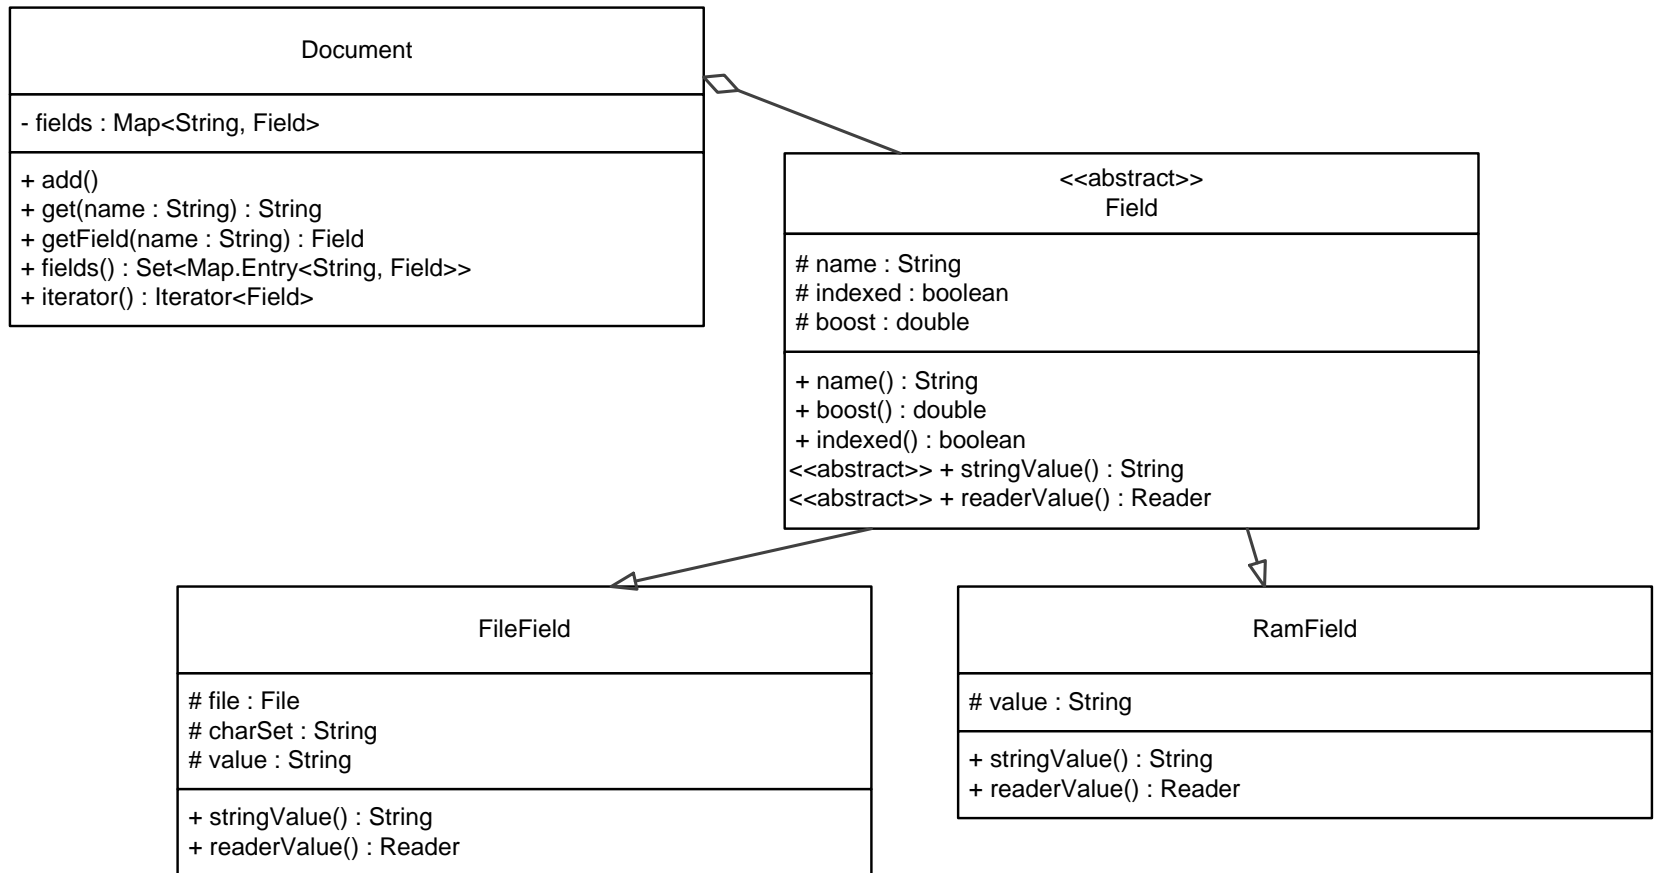
- Checkout the "recommender" project from the repository
- Packages:
 - `document` (for step 1)
 - `analysis` (for step 2)
 - `index and store` (for step 3)
 - `search` (for step 4)

Document (1/2)

- A class that general enough to contain the searchable content
- Can we represent a document as a String object?
- Yes, but
 - Sometimes, some fields (e.g., title) of a document is more important than others (e.g., description)
 - We needs a way to boost terms coming from such fields when calculating the cosine similarity
 - You may not be able to load the whole text into memory a time
 - E.g., when text is loaded from a (large) file, or is transmitted from a remote machine


Document (2/2)

- A more rigorous design:



The Text Analyzer Framework

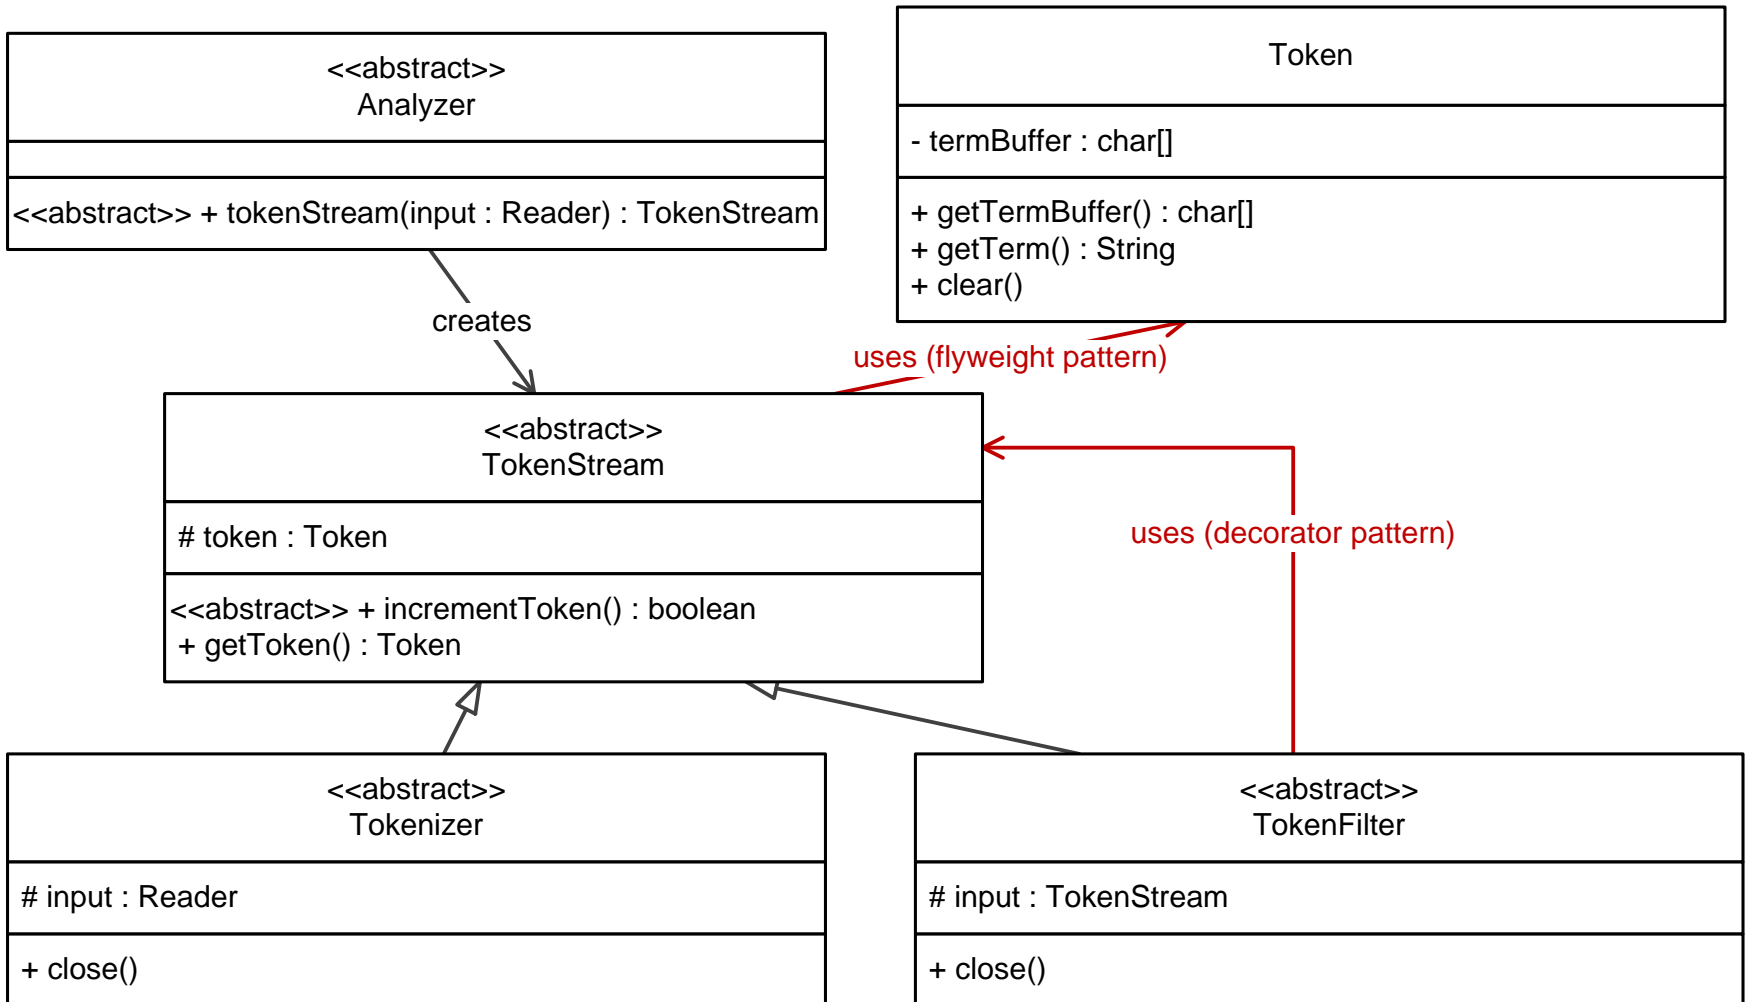
- Input: A text
- Output: terms
- Steps:



Why not String?

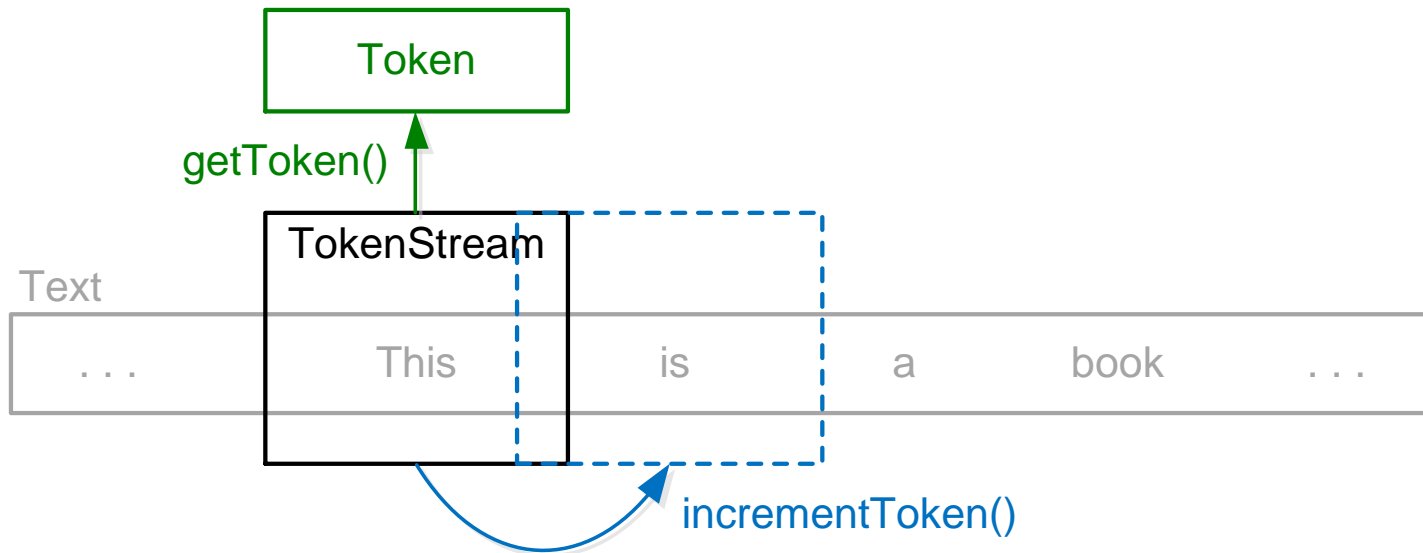
- Recall that String provides a method `split()`?
 - We can break a string `s` in to tokens by `s.split(" ")`
 - More complex split regex can be instructed
- Why do we need to implement an analyzer framework ourselves?
- Again, you may not be able to load the whole text into memory a time
 - E.g., when text is loaded from a (large) file, or is transmitted from a remote machine
- We need a text analyzer framework reads the text (from the beginning to the end) only **once** to parse all the tokens

UML



Core Classes - TokenStream

- Iterates tokens (likes a java.lang.Iterator)
- Call **incrementToken()** (like next()) to move to the next term
- When **getToken()** is called, returns an Token

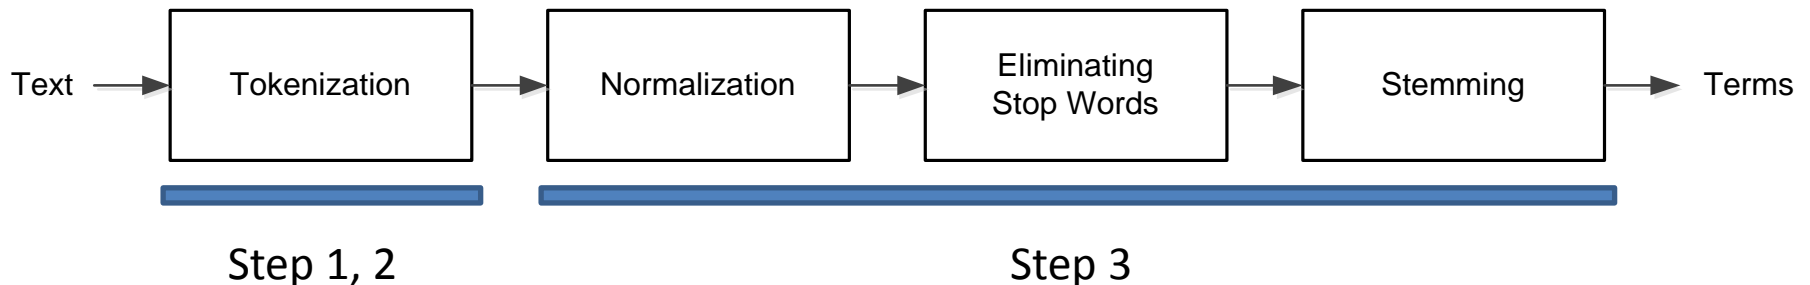


The Flyweight Pattern

- We know that a text can be large
 - There can be many, many tokens
 - This implies that we will create many, many Token instances
- It costs a lot of memory!
- Flyweight pattern
 - Prevents excessive object creations by reusing the same “flyweight” objects
- TokenStream **reuses the same Token instance** to carry different terms during the iteration
 - The getToken() always returns the same instance
 - A care must be taken not to create new objects inside this token instance

Core Classes - TokenStream

- Each time when **incrementToken()** is called:
 1. Reads the beginning chars of the (remaining) text
 2. Assembles chars as a term
 3. Post-processes the term (e.g., lowers case, eliminates it if it is a stop word, performs stemming, etc.). If the term is eliminated, repeats steps 1-3; otherwise, stops reading



Core Classes

- Token
 - Carries the current term
- Tokenizer
 - A TokenStream that performs the steps 1 and 2
 - Works at **char** level
- TokenFilter
 - A TokenStream that performs step 3
 - Works at **word** level

We can implement different Tokenizers and TokenFilters

- Tokenizer
 - E.g. WhiteSpaceTokenizer, LetterTokenizer, etc.
- TokenFilter
 - E.g. LowerCaseTokenFilter, StopFilter, StemFilter, etc.
- The Tokenizers and TokenFilters can be combined arbitrarily
- Remember something we have learnt?

The Decorator Pattern

- Accepts a reference to the preceding TokenStream implementation as the constructor of the succeeding class

```
1. public class StopFilter ... {  
2.     public StopFilter(TokenStream input, ...) {  
3.         this.input = input;  
4.         ...  
5.     }  
6. }
```

- When create the instance, chain the Tokenizers and TokenFilters you want to use in a right order

```
new StemFilter(  
    new StopFilter(  
        new LowerCaseFilter(  
            new LetterTokenizer(input)))
```

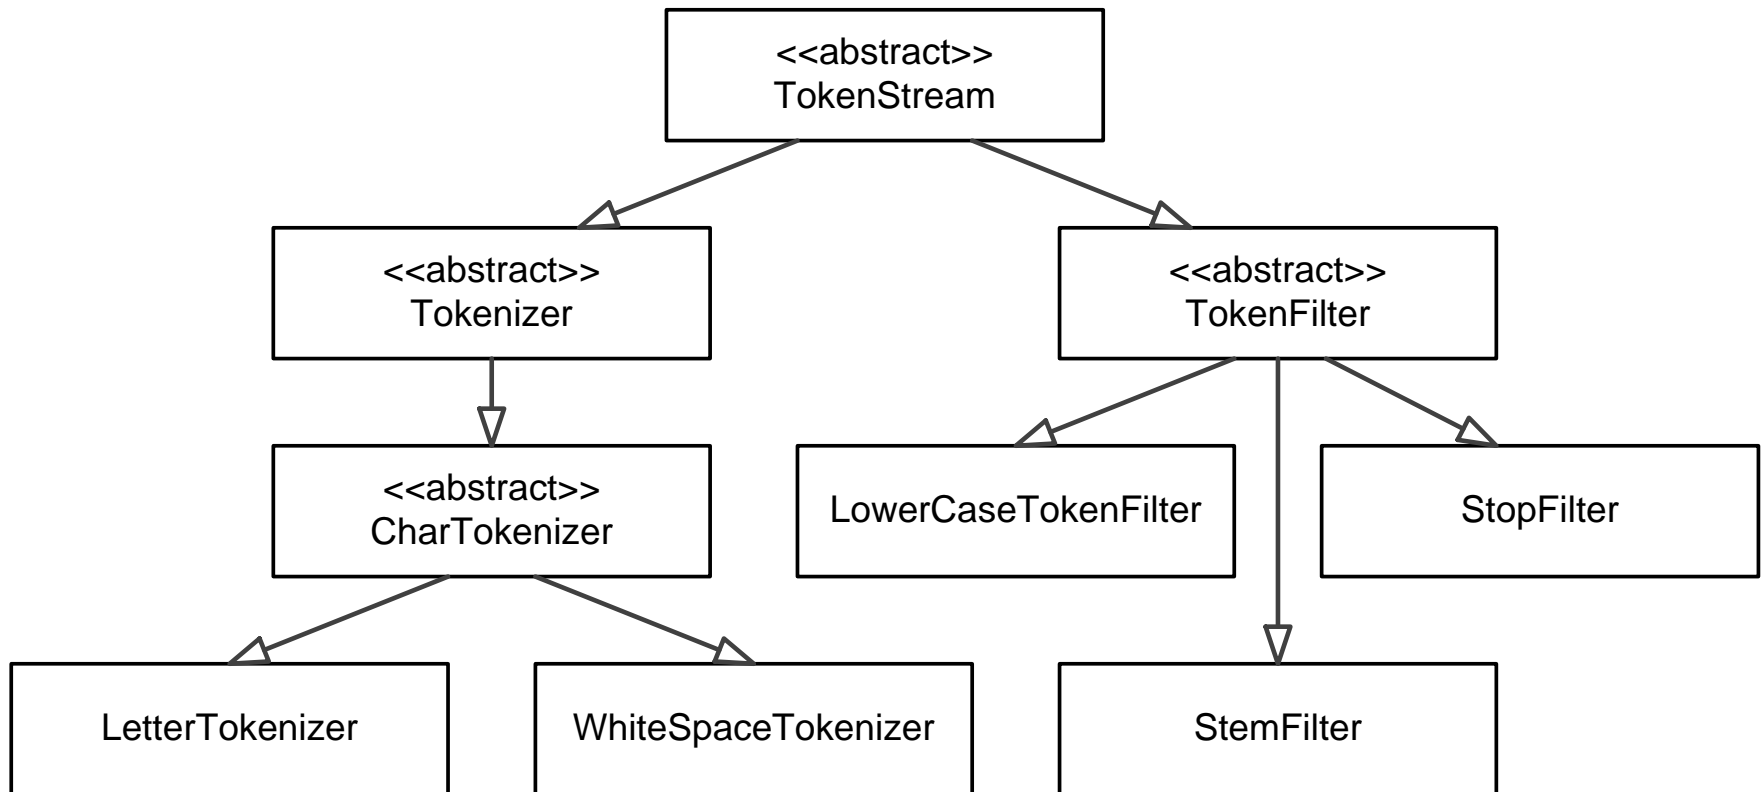
The Decorator Pattern

- When **incrementToken()** is called, every incrementToken() of preceding TokenStream is called. Every TokenStream will complete its mission (e.g. lowercase, remove stop words, etc.) in incrementToken()
- **Decorator pattern**
 - Separates concerns in different TokenStream implementations
 - But **allows dynamic chaining** of these implementations

Core Classes - Analyzer

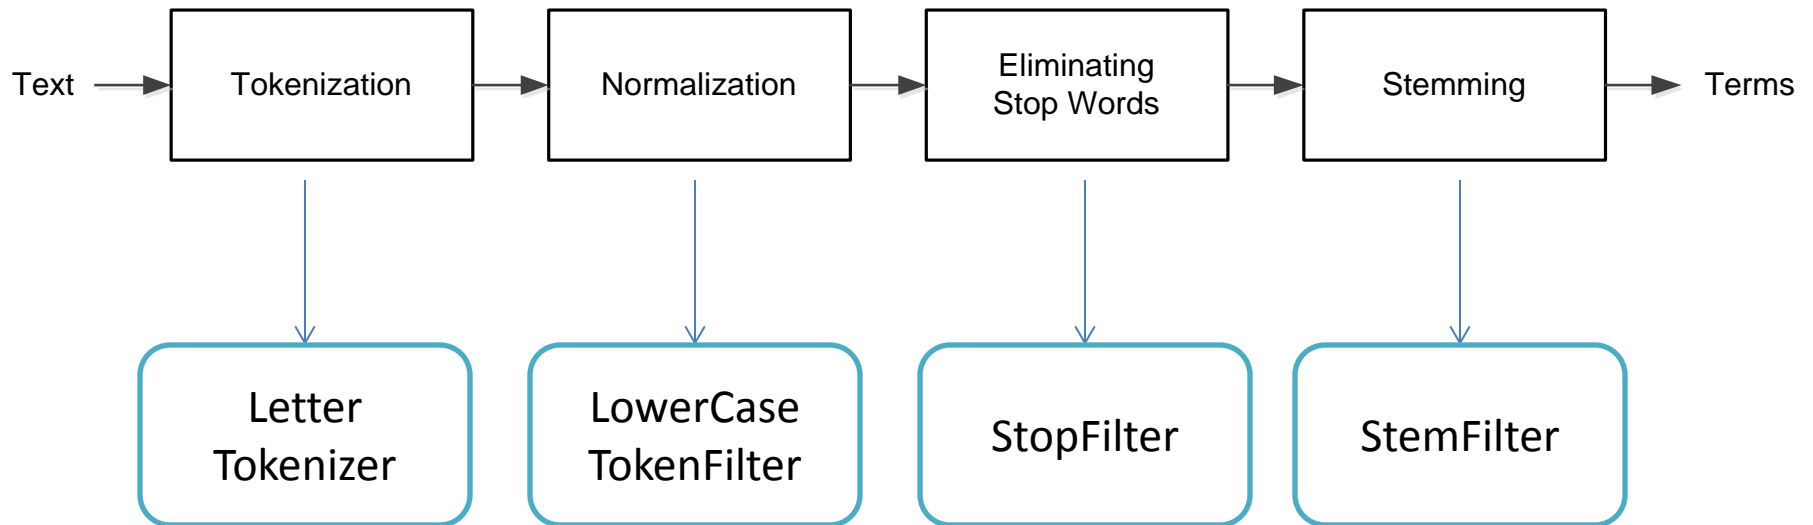
- Accepts the Reader of a text as input
 - A reader can load only the currently reading characters into memory
- It is the Analyzer that does the decorator chaining
 - You can write a specific Analyzer for specific texts without rewriting the TokenStream implementations
 - Add new TokenStream only when you need a new “step” in the analyzing process
 - When `tokenStream()` is called, creates an TokenStream instance (so users can use that instance to iterate tokens)
 - The instance can be “combined” from other TokenStream instances (can be of type Tokenizer or TokenFilter)

TokenStream Hierarchy



The Text Analyzer Framework

- Input: A text
- Output: terms
- Steps:



Today's Mission

- We use the flyweight `Token` object to avoid excessive creation of `String` objects
- However, currently, the `BadLowerCaseFilter` and `BadStopFilter` will create a `String` object during each call to the `incrementToken()`
 - Create excessive `String` objects
- TODO:
 - Implement the `LowerCaseFilter` and `StopFilter` that avoid this problem
 - Let `DefAnalyzer` use the `LowerCaseFilter` and `StopFilter`

Hints

- The main method used today is in `DefAnalyzerApp` in `recommender` package
- In the `LowerCaseFilter`, try get `char[]` from the `Token` object, and make each char lowercase
- In the `StopFilter`, try to use the `CharArraySet.contains()` to determine if the `Token` object contains a stop word